

PA3

March 25, 2021

```
[1]: import numpy as np

[2]: def extract_data(file):
    data = []
    with open(file) as f:
        for lines in f.readlines():
            data.append(np.array([int(i) for i in lines.strip().split(" ")]))

    return data

[3]: def extract_dict(file):
    data = []
    with open(file) as f:
        for lines in f.readlines():
            data.append(lines.strip())

    return data

[4]: # Extracting the data
train = extract_data("pa3train.txt")
test = extract_data("pa3test.txt")

[5]: # Extracting the dictionary
word_dict = extract_dict("pa3dictionary.txt")

[6]: # Question 1:
reduced_train = []
reduced_test = []
weights = np.zeros(len(train[1]) - 1)
output = []
incorrect = 0
total = 0
passes = 1

for vector in train:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_train.append(vector)
```

```

for vector in test:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_test.append(vector)

for i in range(passes):
    for vector in reduced_train:
        if (vector[-1] == 1):
            y = -1
        elif (vector[-1] == 2):
            y = 1

        if (y * np.dot(weights, vector[:-1]) <= 0):
            weights += y*vector[:-1]

for i in range(len(reduced_train)):
    if (reduced_train[i][-1] == 1):
        y = -1
    elif (reduced_train[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_train[i][:-1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("Perceptron Classifier w/ 1 pass")
print("training error: " + str(error))

incorrect = 0
total = 0
for i in range(len(reduced_test)):
    if (reduced_test[i][-1] == 1):
        y = -1
    elif (reduced_test[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_test[i][:-1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

```

```
print("test error: " + str(error))
```

Perceptron Classifier w/ 1 pass
training error: 0.04128440366972477
test error: 0.05305039787798409

```
[7]: # Question 1:
reduced_train = []
reduced_test = []
weights = np.zeros(len(train[1]) - 1)
output = []
incorrect = 0
total = 0
passes = 2

for vector in train:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_train.append(vector)

for vector in test:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_test.append(vector)

for i in range(passes):
    for vector in reduced_train:
        if (vector[-1] == 1):
            y = -1
        elif (vector[-1] == 2):
            y = 1

        if (y * np.dot(weights, vector[:-1]) <= 0):
            weights += y*vector[:-1]

    for i in range(len(reduced_train)):
        if (reduced_train[i][-1] == 1):
            y = -1
        elif (reduced_train[i][-1] == 2):
            y = 1

        output = np.sign(np.dot(weights, reduced_train[i][:-1]))
        total += 1

        if (output != y):
            incorrect += 1

error = incorrect/total
```

```

print("Perceptron Classifier w/ 2 passes")
print("training error: " + str(error))

incorrect = 0
total = 0
for i in range(len(reduced_test)):
    if (reduced_test[i][-1] == 1):
        y = -1
    elif (reduced_test[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_test[i][: -1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("test error: " + str(error))

```

Perceptron Classifier w/ 2 passes
training error: 0.04036697247706422
test error: 0.0610079575596817

```

[8]: # Question 1:
reduced_train = []
reduced_test = []
weights = np.zeros(len(train[1]) - 1)
output = []
incorrect = 0
total = 0
passes = 3

for vector in train:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_train.append(vector)

for vector in test:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_test.append(vector)

for i in range(passes):
    for vector in reduced_train:
        if (vector[-1] == 1):
            y = -1

```

```

        elif (vector[-1] == 2):
            y = 1

        if (y * np.dot(weights, vector[:-1]) <= 0):
            weights += y*vector[:-1]

for i in range(len(reduced_train)):
    if (reduced_train[i][-1] == 1):
        y = -1
    elif (reduced_train[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_train[i][:-1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("Perceptron Classifier w/ 3 passes")
print("training error: " + str(error))

incorrect = 0
total = 0
for i in range(len(reduced_test)):
    if (reduced_test[i][-1] == 1):
        y = -1
    elif (reduced_test[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_test[i][:-1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("test error: " + str(error))

```

Perceptron Classifier w/ 3 passes
 training error: 0.02110091743119266
 test error: 0.04509283819628647

```

[9]: # Question 1:
      reduced_train = []

```

```

reduced_test = []
weights = np.zeros(len(train[1]) - 1)
output = []
incorrect = 0
total = 0
passes = 4

for vector in train:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_train.append(vector)

for vector in test:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_test.append(vector)

for i in range(passes):
    for vector in reduced_train:
        if (vector[-1] == 1):
            y = -1
        elif (vector[-1] == 2):
            y = 1

        if (y * np.dot(weights, vector[:-1]) <= 0):
            weights += y*vector[:-1]

    for i in range(len(reduced_train)):
        if (reduced_train[i][-1] == 1):
            y = -1
        elif (reduced_train[i][-1] == 2):
            y = 1

        output = np.sign(np.dot(weights, reduced_train[i][:-1]))
        total += 1

        if (output != y):
            incorrect += 1

error = incorrect/total

print("Perceptron Classifier w/ 4 passes")
print("training error: " + str(error))

incorrect = 0
total = 0
for i in range(len(reduced_test)):
    if (reduced_test[i][-1] == 1):
        y = -1

```

```

elif (reduced_test[i][-1] == 2):
    y = 1

output = np.sign(np.dot(weights, reduced_test[i][: -1]))
total += 1

if (output != y):
    incorrect += 1

error = incorrect/total

print("test error: " + str(error))

```

Perceptron Classifier w/ 4 passes
training error: 0.01926605504587156
test error: 0.04774535809018567

```

[10]: # Question 2:
reduced_train = []
reduced_test = []
weights = np.zeros(len(train[1]) - 1)
learning_rate = 0.001
iterations = 2
incorrect = 0
total = 0

for vector in train:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_train.append(vector)

for vector in test:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_test.append(vector)

for i in range(iterations):
    loss = []

    for vector in reduced_train:
        if (vector[-1] == 1):
            y = -1
        elif (vector[-1] == 2):
            y = 1

        loss.append((y * vector[: -1]) / (1 + np.exp(y * np.dot(weights, vector[:
→ -1])))))

```

```

weights += learning_rate * np.sum(loss, axis = 0)

for i in range(len(reduced_train)):
    if (reduced_train[i][-1] == 1):
        y = -1
    elif (reduced_train[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_train[i][: -1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("Logistic Regression Classifier w/ 2 iterations")
print("training error: " + str(error))

incorrect = 0
total = 0
for i in range(len(reduced_test)):
    if (reduced_test[i][-1] == 1):
        y = -1
    elif (reduced_test[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_test[i][: -1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("test error: " + str(error))

```

Logistic Regression Classifier w/ 2 iterations
training error: 0.4972477064220184
test error: 0.4907161803713528

```

[11]: # Question 2:
reduced_train = []
reduced_test = []
weights = np.zeros(len(train[1]) - 1)
learning_rate = 0.001
iterations = 10

```



```

incorrect = 0
total = 0

for vector in train:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_train.append(vector)

for vector in test:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_test.append(vector)

for i in range(iterations):
    loss = []

    for vector in reduced_train:
        if (vector[-1] == 1):
            y = -1
        elif (vector[-1] == 2):
            y = 1

        loss.append((y * vector[:-1]) / (1 + np.exp(y * np.dot(weights, vector[:
↪-1]))))

    weights += learning_rate * np.sum(loss, axis = 0)

for i in range(len(reduced_train)):
    if (reduced_train[i][-1] == 1):
        y = -1
    elif (reduced_train[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_train[i][: -1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect / total

print("Logistic Regression Classifier w/ 10 iterations")
print("training error: " + str(error))

incorrect = 0
total = 0
for i in range(len(reduced_test)):
    if (reduced_test[i][-1] == 1):

```

```

        y = -1
    elif (reduced_test[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_test[i][: -1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("test error: " + str(error))

```

Logistic Regression Classifier w/ 10 iterations

training error: 0.29724770642201837

test error: 0.29708222811671087

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:28: RuntimeWarning:
overflow encountered in exp

```

[12]: # Question 2:
reduced_train = []
reduced_test = []
weights = np.zeros(len(train[1]) - 1)
learning_rate = 0.001
iterations = 50
incorrect = 0
total = 0

for vector in train:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_train.append(vector)

for vector in test:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_test.append(vector)

for i in range(iterations):
    loss = []

    for vector in reduced_train:
        if (vector[-1] == 1):
            y = -1
        elif (vector[-1] == 2):
            y = 1

```

```

        loss.append((y * vector[:-1])/(1 + np.exp(y * np.dot(weights, vector[:
→-1]))))

    weights += learning_rate * np.sum(loss, axis = 0)

for i in range(len(reduced_train)):
    if (reduced_train[i][-1] == 1):
        y = -1
    elif (reduced_train[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_train[i][::-1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("Logistic Regression Classifier w/ 50 iterations")
print("training error: " + str(error))

incorrect = 0
total = 0
for i in range(len(reduced_test)):
    if (reduced_test[i][-1] == 1):
        y = -1
    elif (reduced_test[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_test[i][::-1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("test error: " + str(error))

```

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:28: RuntimeWarning:
overflow encountered in exp

Logistic Regression Classifier w/ 50 iterations
training error: 0.03944954128440367
test error: 0.0610079575596817

```

[13]: # Question 2:
reduced_train = []
reduced_test = []
weights = np.zeros(len(train[1]) - 1)
learning_rate = 0.001
iterations = 100
incorrect = 0
total = 0

for vector in train:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_train.append(vector)

for vector in test:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_test.append(vector)

for i in range(iterations):
    loss = []

    for vector in reduced_train:
        if (vector[-1] == 1):
            y = -1
        elif (vector[-1] == 2):
            y = 1

        loss.append((y * vector[:-1]) / (1 + np.exp(y * np.dot(weights, vector[:
↪-1])))))

    weights += learning_rate * np.sum(loss, axis = 0)

for i in range(len(reduced_train)):
    if (reduced_train[i][-1] == 1):
        y = -1
    elif (reduced_train[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_train[i][: -1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect / total

print("Logistic Regression Classifier w/ 100 iterations")

```

```

print("training error: " + str(error))

incorrect = 0
total = 0
for i in range(len(reduced_test)):
    if (reduced_test[i][-1] == 1):
        y = -1
    elif (reduced_test[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_test[i][: -1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("test error: " + str(error))

```

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:28: RuntimeWarning: overflow encountered in exp

Logistic Regression Classifier w/ 100 iterations
training error: 0.02018348623853211
test error: 0.04509283819628647

```

[78]: # Question 3:
reduced_train = []
reduced_test = []
weights = np.zeros(len(train[1]) - 1)
output = []
incorrect = 0
total = 0
passes = 3

for vector in train:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_train.append(vector)

for vector in test:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_test.append(vector)

for i in range(passes):
    for vector in reduced_train:
        if (vector[-1] == 1):
            y = -1

```

```

        elif (vector[-1] == 2):
            y = 1

        if (y * np.dot(weights, vector[:-1]) <= 0):
            weights += y*vector[:-1]

for i in range(len(reduced_train)):
    if (reduced_train[i][-1] == 1):
        y = -1
    elif (reduced_train[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_train[i][:-1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("Perceptron Classifier w/ 3 passes")
print("training error: " + str(error))

incorrect = 0
total = 0
for i in range(len(reduced_test)):
    if (reduced_test[i][-1] == 1):
        y = -1
    elif (reduced_test[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_test[i][:-1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("test error: " + str(error))
print(" ")

smallest_3_values_index = np.argsort(weights)[:3]
smallest_3_values = [weights[i] for i in smallest_3_values_index]

biggest_3_values_index = np.argsort(weights)[-3:][::-1]
biggest_3_values = [weights[i] for i in biggest_3_values_index]

```

```

print("Most positively correlated words: ")
for i in range(len(biggest_3_values_index)):
    print(word_dict[biggest_3_values_index[i]])

print(" ")

print("Most negatively correlated words: ")
for i in range(len(smallest_3_values_index)):
    print(word_dict[smallest_3_values_index[i]])

```

Perceptron Classifier w/ 3 passes
training error: 0.02110091743119266
test error: 0.04509283819628647

Most positively correlated words:
he
team
game

Most negatively correlated words:
file
program
line

```

[81]: # Question 4:
reduced_train = []
reduced_test = []
weights = np.zeros(len(train[1]) - 1)
learning_rate = 0.001
iterations = 50
incorrect = 0
total = 0

for vector in train:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_train.append(vector)

for vector in test:
    if (vector[-1] == 1 or vector[-1] == 2):
        reduced_test.append(vector)

for i in range(iterations):
    loss = []

    for vector in reduced_train:

```

```

        if (vector[-1] == 1):
            y = -1
        elif (vector[-1] == 2):
            y = 1

        loss.append((y * vector[:-1])/(1 + np.exp(y * np.dot(weights, vector[:
↪-1]))))

    weights += learning_rate * np.sum(loss, axis = 0)

for i in range(len(reduced_train)):
    if (reduced_train[i][-1] == 1):
        y = -1
    elif (reduced_train[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_train[i][:-1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("Logistic Regression Classifier w/ 50 iterations")
print("training error: " + str(error))

incorrect = 0
total = 0
for i in range(len(reduced_test)):
    if (reduced_test[i][-1] == 1):
        y = -1
    elif (reduced_test[i][-1] == 2):
        y = 1

    output = np.sign(np.dot(weights, reduced_test[i][:-1]))
    total += 1

    if (output != y):
        incorrect += 1

error = incorrect/total

print("test error: " + str(error))
print(" ")

smallest_3_values_index = np.argsort(weights)[:3]

```



```

smallest_3_values = [weights[i] for i in smallest_3_values_index]

biggest_3_values_index = np.argsort(weights)[-3:][::-1]
biggest_3_values = [weights[i] for i in biggest_3_values_index]

print("Most positively correlated words: ")
for i in range(len(biggest_3_values_index)):
    print(word_dict[biggest_3_values_index[i]])

print(" ")

print("Most negatively correlated words: ")
for i in range(len(smallest_3_values_index)):
    print(word_dict[smallest_3_values_index[i]])

```

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:28: RuntimeWarning:
overflow encountered in exp

Logistic Regression Classifier w/ 50 iterations
training error: 0.03944954128440367
test error: 0.0610079575596817

Most positively correlated words:

he
game
they

Most negatively correlated words:

window
file
use

```

[180]: # Question 5:
labels = [1, 2, 3, 4, 5, 6]
C = np.zeros((7,6))
num_labels = np.zeros(6)
weights_matrix = []

for label in labels:
    weights = np.zeros(len(train[1]) - 1)

    for vector in train:
        if (vector[-1] == label):
            y = 1
            num_labels[label - 1] += 1
        elif (vector[-1] != label):
            y = -1

```

```

        if (y * np.dot(weights, vector[:-1]) <= 0):
            weights += y*vector[:-1]

weights_matrix.append(weights)

for i in range(len(train)):
    output = np.sign(np.dot(weights_matrix, train[i][:-1]))
    s = np.sort(output)

    if (s[-1] == s[-2]):
        C[6][train[i][-1] - 1] += 1
    elif (np.argmax(output) + 1 != train[i][-1]):
        C[np.argmax(output)][train[i][-1] - 1] += 1
    elif (np.argmax(output) + 1 == train[i][-1]):
        C[np.argmax(output)][train[i][-1] - 1] += 1

C = C/num_labels
print("Confusion Matrix for training data: ")
print(C)

print(" ")
print("Confusion Matrix for test data: ")

C = np.zeros((7,6))
num_labels = np.zeros(6)

for label in labels:
    for vector in test:
        if(vector[-1] == label):
            num_labels[label - 1] += 1

for i in range(len(test)):
    output = np.sign(np.dot(weights_matrix, test[i][:-1]))
    s = np.sort(output)

    if (s[-1] == s[-2]):
        C[6][test[i][-1] - 1] += 1
    elif (np.argmax(output) + 1 != test[i][-1]):
        C[np.argmax(output)][test[i][-1] - 1] += 1
    elif (np.argmax(output) + 1 == test[i][-1]):
        C[np.argmax(output)][test[i][-1] - 1] += 1

C = C/num_labels
print(C)

```

```

# (a)
# C5 is the perceptron classifier w/ highest accuracy.
# This means that when given a feature vector from the test dataset w/ label 5,
  ↳ it correctly classifies w/ an accuracy of 80%

# (b)
# C3 is the perceptron classifier w/ lowest accuracy.
# This means that when given a feature vector from the test dataset w/ label 3,
  ↳ it correctly classifies w/ an accuracy of 37%

# (c)
# Looking at the Confusion Matrix for the test dataset, the perceptron
  ↳ classifier had the most mistakes when
# trying to classify feature vectors w/ label 6, and mostly outputted labels 5
  ↳ instead.
# i = 5 j = 6

```

Confusion Matrix for training data:

```

[[0.80110497 0.00365631 0.00956023 0.00386847 0.0019305  0.00284091]
 [0.01104972 0.7678245  0.03441683 0.01740812 0.01351351 0.01988636]
 [0.          0.00182815 0.43403442 0.00580271 0.00579151 0.00852273]
 [0.01473297 0.01462523 0.01529637 0.71760155 0.00772201 0.00852273]
 [0.0092081  0.02010969 0.03441683 0.00773694 0.78957529 0.10511364]
 [0.00368324 0.00914077 0.02294455 0.00386847 0.03474903 0.50568182]
 [0.16022099 0.18281536 0.44933078 0.24371373 0.14671815 0.34943182]]

```

Confusion Matrix for test data:

```

[[0.71891892 0.01041667 0.03428571 0.02173913 0.          0.          ]
 [0.01081081 0.65625     0.03428571 0.02717391 0.01282051 0.01851852]
 [0.          0.015625    0.37142857 0.          0.          0.02777778]
 [0.01621622 0.00520833 0.          0.69021739 0.          0.          ]
 [0.01621622 0.03125     0.07428571 0.00543478 0.80128205 0.12037037]
 [0.00540541 0.01041667 0.03428571 0.          0.07051282 0.5          ]
 [0.23243243 0.27083333 0.45142857 0.25543478 0.11538462 0.33333333]]

```

[]: